

# **Secure Distributed Data and Event Processing at Scale**

## **Where are we now?**

**P. Eugster USI Lugano**

**Jun 30 DEBS'23**

# Roadmap

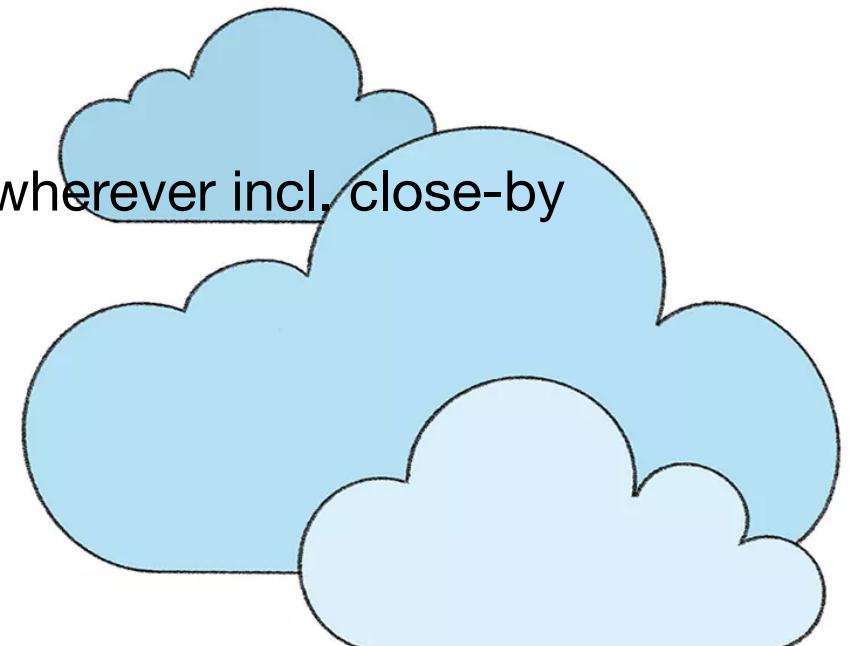
- Motivation: distributed data and event processing
- Problem: untrusted third-party infrastructure
- Solution fragments: security mechanisms
- Requirements: what do we need
- Tradeoffs: what comes first
- HYDRA: a carefully balanced approach
- Conclusions: lessons learned and open challenges



# Motivation

## Distributed data and event processing

- Rise of “remote execution”
  - Advent of cloud computing, edge
  - Massive resources available on demand, wherever incl. close-by
  - “Distribution by commodity”
- Further digitization of society
  - IoT, wearable computing, ...
  - “Distribution by necessity”



# Problem

## Untrusted third-party infrastructure

- So you want to do distributed processing on large data sets?
  - Do not have your own cluster?
  - Just use cloud or edge data centers, or both!
  - And say goodbye to your data...
- Causes
  - Multitenancy, internal threats, big government
  - Use of “foreign” software, tool sprawl
- Almost half of data breaches happen in cloud, with average cost of \$4.35M!!! [IBM’22]

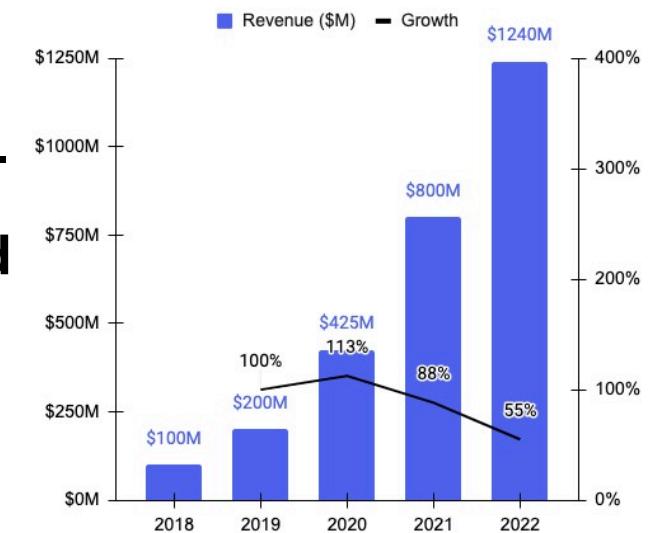


[IBM’22] <https://www.ibm.com/reports/data-breach>

# Data Analytics and Confidential Computing

## Relevance for society and economy

- 2022 revenue of DataBricks, a leader in data analytics, estimated at \$1.24B up from \$800M in 2021, 90% from platform based on Spark [Sabra'23].
- DataBricks/Spark lacks support for security **required** by many scenarios.
- Gartner predicts size of **confidential computing** market to be \$54B by 2026 [Gartner'23].



[Sabra'23] <https://sacra.com/c/databricks/>

[Gartner'23] <https://www.gartner.com/en/newsroom/press-releases/2021-03-23-gartner-identifies-top-security-and-risk-management-t>

# Solution Fragments

A word cloud diagram centered around the word "homomorphic". The word is in a large, bold, purple font. Surrounding it are various other words in different colors and sizes, representing different solution fragments. The words include:

- Encryption (large, central, purple)
- Functional (yellow)
- Sgx (yellow)
- Secure (purple)
- Memory (blue)
- Computing (red)
- Garbled (red)
- Trusted (orange)
- Circuit (red)
- Extensions (blue)
- Sme (red)
- Multi-party (red)
- Software (blue)
- Attribute-based (orange)
- Enclave (red)
- Nitro (blue)
- Partially (yellow)
- Tee (blue)
- Guard (blue)
- Fully (yellow)
- Environment (orange)
- Execution (orange)
- Fhe (yellow)
- Phe (yellow)

# Mechanisms

## Software (SW)

- Cryptographic primitives, e.g.,
  - Fully homomorphic encryption [Gentry'10]
  - Partially homomorphic encryption (PHE)
  - Functional encryption [Boneh et al.'11]
  - Attributed-based encryption [Shamir'84]
  - Garbled circuits [Yao'86]
  - ...

[Gentry'10] C. Gentry. Computing Arbitrary Functions of Encrypted Data. Communications of the ACM 53(3): 97-105, 2010.

[Boneh et al.'11] D. Boneh, A. Sahai, and B. Waters: Functional Encryption: Definition and Challenges. TCC 2011.

[Shamir'84] A. Shamir: Identity-Based Cryptosystems and Signature Schemes. Advances in Cryptology 1984.

[Yao'86] A. C.-C. Yao: How to Generate and Exchange Secrets. FOCS 1986.

# Mechanisms

## Hardware (HW)

- Trusted execution environments, e.g.,
  - Intel Secure Guard eXtensions (SGX)
  - AMD Secure Memory Encryption (SME)
  - ARM Trustzone
  - Amazon Nitro
  - ...

# Mechanisms

## Hardware (HW)

- Trusted execution environments, e.g.,
  - Intel Secure Guard eXtensions (SGX)
  - AMD Secure Memory Encryption (SME)
  - ARM Trustzone
  - Amazon Nitro
  - ...



**Which one(s) to use? How?**

# Requirements?

**More/higher/stronger is better**

- Transparency
  - Very few data analysts are experts in security
- Efficiency
  - Many potential users deterred by high overhead w.r.t. security-agnostic execution
  - 10x slowdown critical
- Portability
  - Lock-in with cloud or platform providers is never good, may support different (HW) mechanisms
  - New mechanisms keep being proposed

# Requirements (cont'd)?

## Higher/stronger is better

- Interoperability
  - Important especially if combining cloud and edge in *data pipelines*, notably continuous processing
  - Interoperability  $\not\Rightarrow$  portability: if components can interact across platforms doesn't mean they can move
  - For simplicity we consider interoperability as platform-independence including portability
- Security
  - What is the attacker/threat model and guarantees given in the face of those? If little security is offered ...?
  - How strongly are they substantiated? Informal, "formal" on paper, mechanically verified, ...?
- Expressivity
  - Supporting simpler computations is always easier

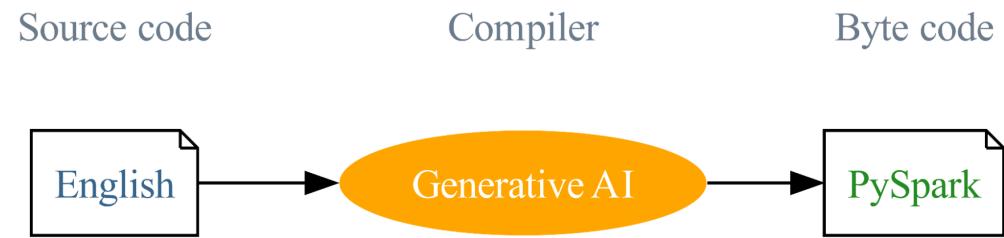
# Tradeoffs

- Transparency vs interoperability: user-facing language?
  - SecureScala [Hauck et al.'16]  
[Hauck et al.'16] M. Hauck, S. Savvides, P. Eugster, M. Mezini, and G. Salvaneschi: SecureScala: Scala Embedding of Securecomputations. SCALA 2016.
  - Transparency vs efficiency: transparency for whom?
    - Cuttlefish [Savvides et al.'17]  
[Savvides et al.'17] S. Savvides, J. J. Stephen, M. Saeida Ardekani, V. Sundaram, and P. Eugster: Secure Data Types: A Simple Abstraction for Confidentiality-Preserving Data Analytics. SoCC 2017.
    - Security vs efficiency: what say you (user)?
      - Symmetria [Savvides et al.'20]  
[Savvides et al.'20] S. Savvides, D. Khandelwal, and P. Eugster: Efficient Confidentiality-Preserving Data Analytics over Symmetrically Encrypted Datasets. VLDB 2020.
      - Expressivity vs efficiency: continuous is always harder
        - STYX [Stephen et al.'16], C3PO [Savvides et al.'22]  
[Stephen et al.'16] J. J. Stephen, S. Savvides, V. Sundaram, M. Saeida Ardekani, and P. Eugster: STYX: Stream Processing with Trustworthy Cloud-based Execution. SoCC 2016.
        - [Savvides et al.'22] S. Savvides, S. Kumar, J. J. Stephen, and P. Eugster: C3PO. Cloud-based Confidentiality-preserving Continuous Query Processing. ACM TOPS 25(1), 2022.

# Tradeoffs

- Transparency vs interoperability: user-facing language?
  - SecureScala [Hauck et al.'16]
- Transparency vs efficiency: transparency for
  - Cuttlefish [Savvides et al.'17]
- Security vs efficiency: what say you (user)?
  - Symmetria [Savvides et al.'20]
- Expressivity vs efficiency: continuous is always harder
  - STYX [Stephen et al.'16], C3PO [Savvides et al.'22]

Databricks introducing English as the new programming language for Apache Spark



[Stephen et al.'16] J. J. Stephen, S. Savvides, V. Sundaram, M. Saeida Ardekani, and P. Eugster: STYX: Stream Processing with Trustworthy Cloud-based Execution. SoCC 2016.

[Savvides et al.'22] S. Savvides, S. Kumar, J. J. Stephen, and P. Eugster: C3PO. Cloud-based Confidentiality-preserving Continuous Query Processing. ACM TOPS 25(1), 2022.

# Tradeoffs

- Transparency vs interoperability: user-facing language?
  - SecureScala [Hauck et al.'16]
- Transparency vs efficiency: transparency for whom?
  - Cuttlefish [Savvides et al.'17]
- Security vs efficiency: what say you (user)?
  - Symmetria [Savvides et al.'20]
- Expressivity vs efficiency: continuous is always harder
  - STYX [Stephen et al.'16], C3PO [Savvides et al.'22]

[Hauck et al.'16] M. Hauck, S. Savvides, P. Eugster, M. Mezini, and G. Salvaneschi: SecureScala: Scala Embedding of Securecomputations. SCALA 2016.

[Savvides et al.'17] S. Savvides, J. J. Stephen, M. Saeida Ardekani, V. Sundaram, and P. Eugster: Secure Data Types: A Simple Abstraction for Confidentiality-Preserving Data Analytics. SoCC 2017.

[Savvides et al.'20] S. Savvides, D. Khandelwal, and P. Eugster: Efficient Confidentiality-Preserving Data Analytics over Symmetrically Encrypted Datasets. VLDB 2020.

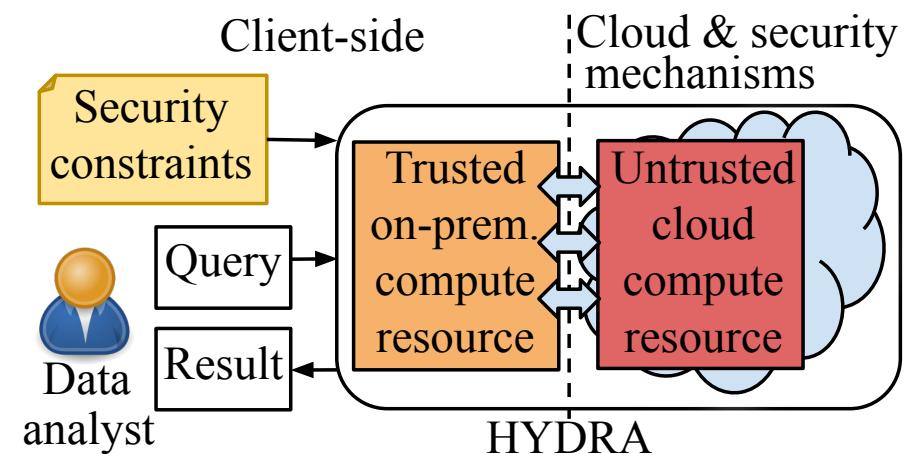
[Stephen et al.'16] J. J. Stephen, S. Savvides, V. Sundaram, M. Saeida Ardekani, and P. Eugster: STYX: Stream Processing with Trustworthy Cloud-based Execution. SoCC 2016.

[Savvides et al.'22] S. Savvides, S. Kumar, J. J. Stephen, and P. Eugster: C3PO. Cloud-based Confidentiality-preserving Continuous Query Processing. ACM TOPS 25(1), 2022.

# HYDRA

## Hybrid approach to confidentiality-preserving data analytics

- Transparency: data analysts write queries *agnostically to security constraints*.
- Efficiency: over *10x faster* than state of the art [Mangipudi et al.'23].
- Interoperability: *combining* software mechanisms (homomorphic encryption schemes, e.g., Paillier [Paillier'99]) and hardware mechanisms (trusted execution environments, e.g., Intel SGX) from *extensible sets*.
- Security: *secure end-to-end* with *formally verified* guarantees.
- (Expressivity: focus on distributed data analytics)



[Paillier'99] P. Paillier: Public-Key Cryptosystems Based on Composite Degree Residuosity Classes. EUROCRYPT 1999.

[Mangipudi et al.'23] S. Mangipudi, P. Chuprikov, P. Eugster, M. Viering and S. Savvides: Generalized Policy-based Non-interference for Efficient Confidentiality-Preservation. PLDI 2023.

# Transparency

- Full transparency is unrealistic, undesired
- **Can define/split different roles (`SecureScala`, `Cuttlefish`)**
  1. *Data analysts*: know queries; but not computer scientists, not expert programmers
    - Should be able to write queries agnostically to security
  2. *Security experts*: know security mechanisms, attacks, cloud platforms, etc.
    - Define (static) security constraints and policies
  3. *Data managers*: know data semantics, understand security basics
    - Define (static) security requirements for data together with 2.

# Efficiency

- SW can outperform HW (Cuttlefish)
  - E.g., PHE > SGX sometimes
  - **Want SW and HW**
- Differences between SW (Symmetria)
- Researchers frantically working on new SW solutions
- Researchers and engineers frantically working on new HW solutions
- Want to be able to integrate *new* solutions without complete redesign

Operation	PHE scheme	Plaintext	PHE	SGX
filter ( $_ = _$ )	AES-ECB [Daemen and Rijmen 2002]	5.4 s	5.8 s	8.0 s
filter range	OPE [Boldyreva et al. 2009]	5.7 s	8.7 s	10.7 s
filter match	SWP [Song et al. 2000]	5.4 s	8.8 s	9.4 s
groupby	AES-ECB [Daemen and Rijmen 2002]	6.2 s	7.2 s	57.7 s
sort	OPE [Boldyreva et al. 2009]	7.2 s	13.0 s	41.0 s
select ( $_ + _$ )	Paillier [Paillier 1999]	5.2 s	19.4 s	7.8 s
select ( $_ \times _$ )	ElGamal [ElGamal 1985]	5.1 s	22.4 s	7.8 s

# Interoperability

- Support **extensible sets of HW and SW mechanisms**
  - Domains  $\mathcal{D}$ : capture HW mechanisms and infrastructure providers
    - Can treat same HW mechanism differently in different infrastructures
  - Crypto(graphic) schemes  $\mathcal{S}$ : capture SW mechanisms
- Why not just 1 dimension/set of mechanisms?
  - Isn't SGX just AES(-NI) in HW?
    - Different characteristics and use
  - Benefits of combining HW with SW mechanisms for *same* data

# Security

- Extensible sets  $\mathcal{D}$  and  $\mathcal{S}$
- How to integrate mechanisms with different guarantees?
  - E.g., different threat models, known attacks [Fei et al.'21]
- Basis: honest-but-curious (HbC) attacker model
  - Focus on confidentiality
- **Let security experts define what they trust and where (Cuttlefish)**
- **Leverage formal techniques to verify guarantees end-to-end (SecureScala)**

[Fei et al.'21] S. Fei, Z. Yan, W. Ding, and H. Xie: Security Vulnerabilities of SGX and Countermeasures: A Survey. ACM CSUR 54(6), 2021.

# Security Policy $\mathbb{S}$

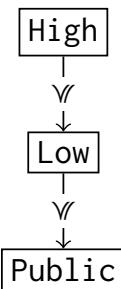
- Custom finite set of labels/levels  $\mathcal{L}$  with lowest element  $\perp$  [Denning'76]
- Arranged according to lattice  $\langle \mathcal{L}, \leq, \sqcup, \sqcap \rangle$ 
  - Partial order  $\leq$  reflexive antisymmetric binary relation on  $\mathcal{L}$ 
    - $l_1 \leq l_2$ : confidentiality requirements imposed by  $l_2$  are at least as strict as  $l_1$ ; always safe to replace  $l_1$  with  $l_2$  when labelling data
  - Least upper bound operator  $\sqcup : \mathcal{L} \times \mathcal{L} \rightarrow \mathcal{L}$
  - Greatest lower bound operator  $\sqcap : \mathcal{L} \times \mathcal{L} \rightarrow \mathcal{L}$
- Security policy  $\mathbb{S}$  assigning levels to combinations of SW and HW mechanisms  $\mathbb{S} : \mathcal{L} \rightarrow 2^{\mathcal{D} \times \mathcal{S}_\emptyset}$

[Denning'76] D. E. Denning: A Lattice Model of Secure Information Flow. Communications of the ACM, 1976.

# Example Security Policy Setup

All steps performed once only

## 1. Define secrecy levels



Cf. US DoD

## 2. Assign mechanisms to levels

Label	Domain	Scheme
$\mathcal{L}$	$\mathcal{D}$	$\mathcal{S}_{\emptyset}$
High	CLNT, SGX	$\emptyset$
High	CLD	AES-GCM
Low	CLD	SWP, AES-ECB, Paillier, ElGamal, OPE

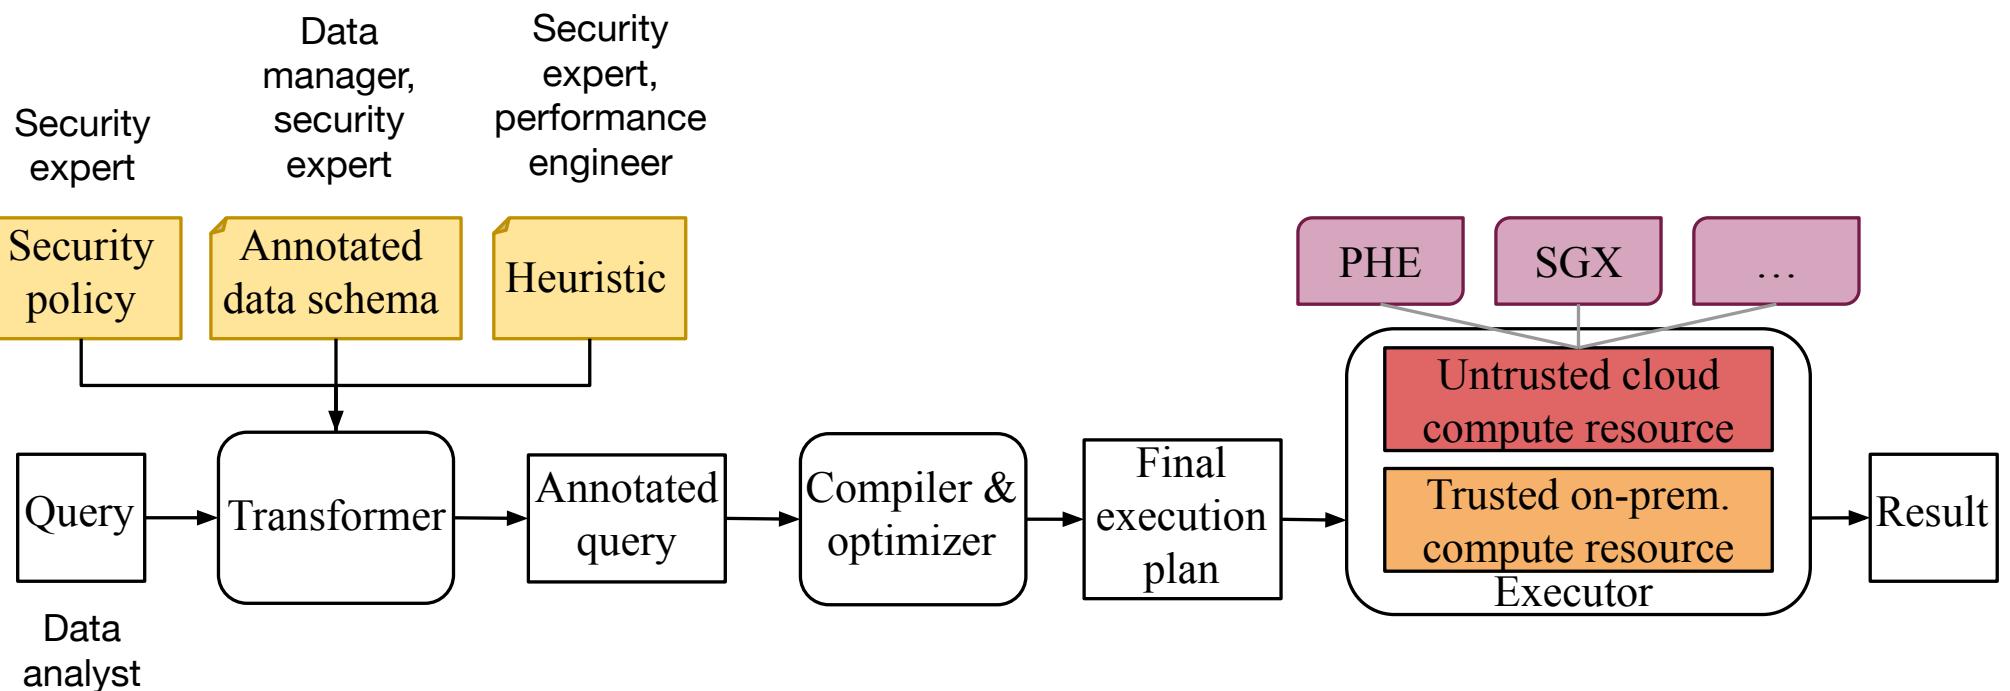
E.g., data at level High can be handled only on the client side (CLNT), by Intel SGX (SGX), or in the cloud (no SGX) only when using AES-GCM encryption

## 3. Label data with levels

Relation	Field	Type	Label	Security type (inferred)
Customers	custId	Str	Low	(Str <sup>AES-ECB</sup> ,Low)
	bal	Dbl	High	(Dbl <sup>AES-GCM</sup> ,High)
	...			
Orders	orderId	Str	High	(Str <sup>AES-GCM</sup> ,High)
	custKey	Str	Low	(Str <sup>AES-ECB</sup> ,Low)
	price	Str	High	(Dbl <sup>AES-GCM</sup> ,High)
	date	Int	Low	(Int <sup>OPE</sup> ,Low)
...				

Security types automatically inferred at query transformation, e.g. AES-GCM encryption for price tags in Orders

# Complete Workflow



# Core Language

## Lambda with relations and relational operators

Type  $\kappa ::= \overline{\kappa} \rightarrow_d \kappa \mid \{f : (p^s, l)\} \mid T\{f : (p^s, l)\} \mid (p^s, l)$

Prim type  $p ::= \text{Int} \mid \text{Dbl} \mid \text{Str} \mid \text{Bool} \mid \dots$

Scheme  $s ::= \emptyset \mid \text{AES-GCM} \mid \text{ElGamal} \mid \text{Paillier} \mid \dots$  (in one-to-one correspondence with  $\mathcal{S}_\emptyset$ )

Domain  $d ::= \text{CLNT} \mid \text{SGX} \mid \text{SEV} \mid \text{CLD} \mid \dots$  (in one-to-one correspondence with  $\mathcal{D}$ )

Value  $v ::= T\{\overline{f : \bar{v}}\} \mid \{\overline{f : v}\} \mid c^s \mid \lambda[d](\overline{x : \kappa}). e \mid f$

Expression  $e ::= v \mid x \mid e(\bar{e}) \mid \oplus(\bar{e}) \mid \{f : e\} \mid e.f \mid \text{table}(name) \mid \theta(\bar{e}) \mid \text{encr}(e, s) \mid \text{decr}(e) \mid [e]_d$

Prim ops  $\oplus ::= + \mid - \mid \dots \mid \wedge \mid \vee \mid \dots$

Query ops  $\theta ::= \text{filter} \mid \text{proj} \mid \text{cross} \mid \text{agg} \mid \dots$

- Blue parts not used by data analysts
- Green parts only used during evaluation

# Transformation

## 1. Original security-agnostic query

---

```
1 agg(filter(cross(table(Customers),
2                     filter(table(Orders),
3                         λ(r0: {/* Orders */}). r0.date < 16052002)),
4                     λ(rCO: {/* Customers + Orders */}). rCO.custId == rCO.custKey),
5                     custId, 0,
6                     λ(rP: {price: Dbl}, acc: Int).
7                     acc + rP.price)
```

---

# Transformation

## 2. Automatically transformed query with security annotations

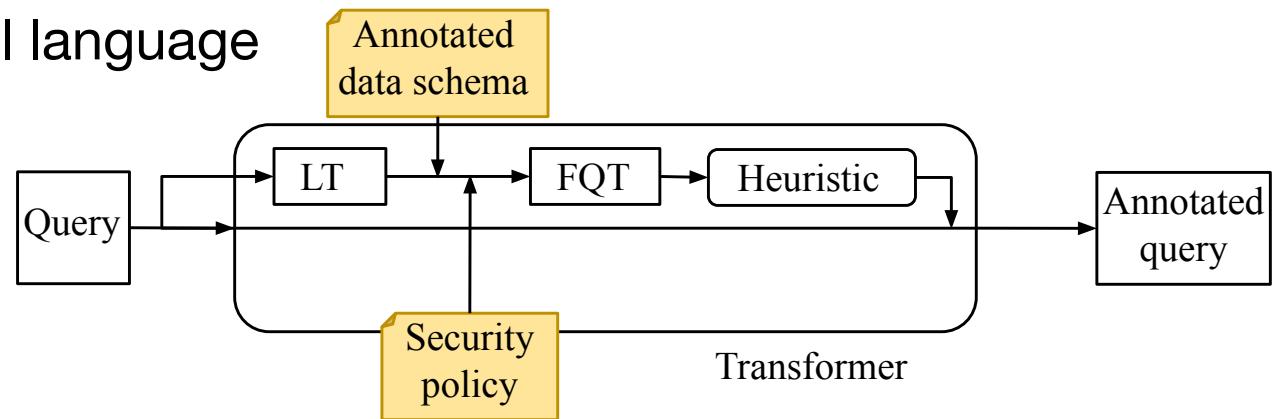
---

```
1 agg(filter(cross(table(Customers),
2                     filter(table(Orders),
3                         λ[SGX](r0: /* Orders */). r0.date < 0x..OPE)),
4                     λ[SGX](rC0: /* Customers + Orders */). rC0.custId == rC0.custKey),
5                     custId, 0x..AES-GCM,
6                     λ[SGX](rP: {price: (DblAES-GCM, High)}, acc: (DblAES-GCM, High)).
7                     encr(decr(acc) + decr(rP.price), AES-GCM))
```

---

# Transformation Heuristics and Framework

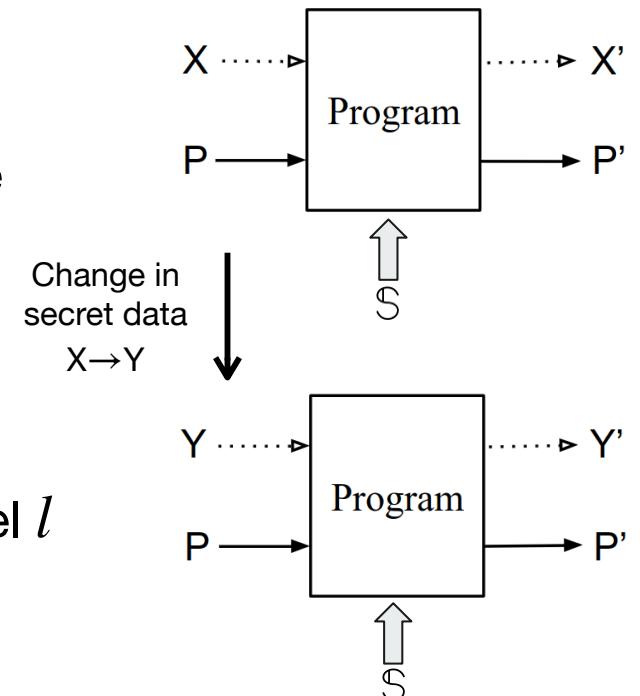
- Formal framework denoting well-formed transformations
  - (Partially) enforced through APIs and libraries for writing heuristics
- Information flow type system verifies transformed queries
  - Also queries written in full language



# Guarantees

## Informal

- Based on *noninterference* [Goguen&Meseguer'82]
  - Generalized to our security policy, thus  $\mathbb{S}$ -noninterference
- Partially ordered set of adversaries  $\langle \mathcal{A}, \subseteq \rangle$ 
  - $A \in \mathcal{A}$  downward-closed set  $A \subseteq \mathcal{D} \times \mathcal{S}_\emptyset$  it can break
  - If  $\mathbb{S}(l) \cap A = \emptyset$  then  $A$  is unable to see any inputs at level  $l$
- Adversary who can't break anything at level  $l$  can't see difference in outputs of computation if inputs at level  $l$  or higher are changed



[Goguen&Meseguer'82] A. Goguen and J. Meseguer:  
Security Policies and Security Models. S&P 1982.

# Guarantees

## Formal

( $\Omega$  table store,  $\rho$  plaintext schema,  $\kappa$  type w/ annotations)

DEFINITION 4 (LEVEL- $l$   $\mathbb{S}$ -NONINTERFERENCE  $\mathbb{S}\text{-NI}(e)_{\rho,d,l}$ ). *Expression  $e$  has  $\mathbb{S}\text{-NI}(e)_{\rho,d,l}$  property dubbed level- $l$   $\mathbb{S}$ -noninterference if and only if for any two stores  $\Omega_1$  and  $\Omega_2$  satisfying  $\rho$ ,  $\Omega_1 \sim_{\rho}^l \Omega_2$ , and any two values  $v_1$  and  $v_2$ ,  $e \xrightarrow[\Omega_1]^* v_1$  and  $e \xrightarrow[\Omega_2]^* v_2$ , it holds that  $v_1 \sim^{d,l} v_2$ .*

DEFINITION 5 ( $\mathbb{S}$ -NONINTERFERENCE  $\mathbb{S}\text{-NI}(e)_{\rho,d}$ ). *Expression  $e$  has  $\mathbb{S}$ -noninterference property  $\mathbb{S}\text{-NI}(e)_{\rho,d}$  if and only if it has level  $l$   $\mathbb{S}$ -noninterference property  $\mathbb{S}\text{-NI}(e)_{\rho,d,l}$  for every  $l$  in  $\mathcal{L}$ .*

THEOREM 1 (SOUNDNESS). *If there exists non-function  $\kappa$ , s.t.,  $\rho \vdash_d e : \kappa$  w.r.t.  $\mathbb{S}$  then  $\mathbb{S}\text{-NI}(e)_{\rho,d}$ .*

THEOREM 2 (SUBJECT REDUCTION). *Let  $\Vdash_{/d} \Omega : \rho, \rho \wr \Gamma \Vdash_d e : \kappa$  and  $e \xrightarrow[\Omega]{} e'$  then  $\rho \wr \Gamma \Vdash_d e' : \kappa$ .*

# Guarantees

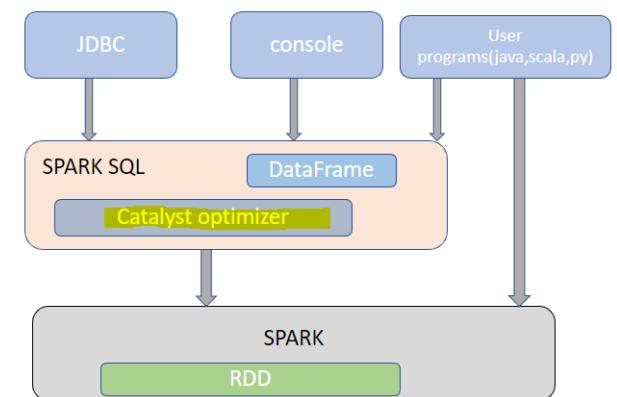
## Transformation

DEFINITION 1 (QUERY TRANSFORMATION). *Function  $\tau[\![\cdot, \cdot]\!]$  is a query transformation iff  $\forall \rho, e$ , and  $(\rho', e') = \tau[\![\rho, e]\!]$ , we have  $e \rightsquigarrow e'$  and  $\rho \rightsquigarrow_S \rho'$ .*

THEOREM 3 (TRANSFORMATION CORRECTNESS). *For query transformation  $\tau[\![\cdot, \cdot]\!]$ , schemata  $\rho$ , and expression  $e$ , let  $(\rho', e') = \tau[\![\rho, e]\!]$ . If  $\rho' \vdash_d e' : \kappa$  for some domain  $d$  and type  $\kappa$ , and also  $e \xrightarrow[\Omega]{}^* v$  for some table store  $\Omega$  satisfying  $\rho$ , then for any  $\Omega' = \text{encrVal}(\Omega, \rho')$ , there exists  $v'$ , s.t.,  $e' \xrightarrow[\Omega']{}^* v'$  and  $\text{decrVal}(v') = v$ .*

# Implementation

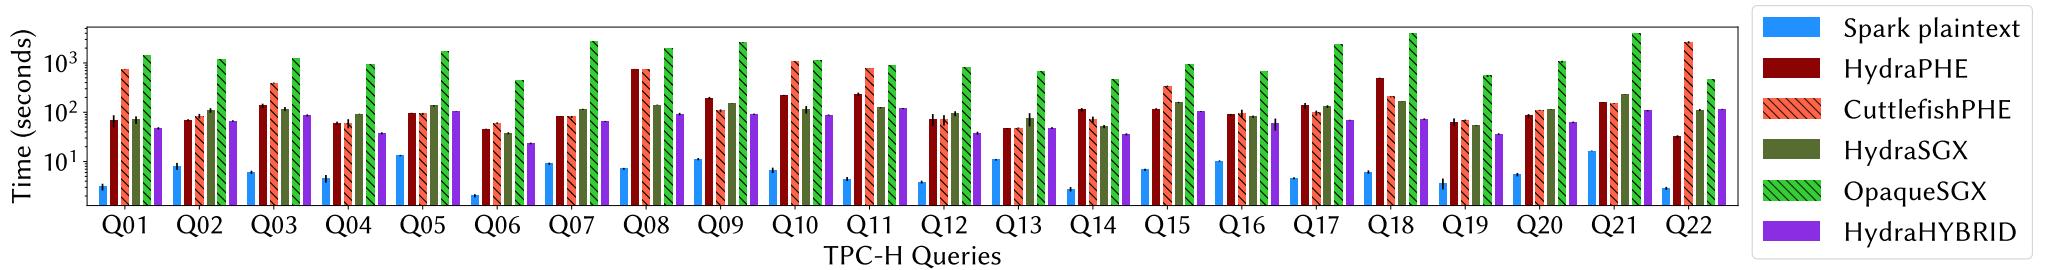
- Built on Apache Spark
  - Used by data analysts through original APIs (SparkSQL)
- Leveraging *Catalyst* extensible query optimizer (query analyzer, optimizer, execution planner)
  - Input data encrypted in analysis phase (3.2 kLoC Scala)
  - Logical optimization in analysis phase (3.4 kLoC Scala)
  - Physical opt. in planning phase (6.2 kLoC Scala and 100 LoC Java)
  - Code generation step (29 kLoC C++, 2.5 kLoC C, 568 LoC Scala)
- SGX used via JNI (mini-interpreter implemented in C)
- Simple HYDRA Hybrid PHE&SGX heuristic: PHE until hitting limit of PHE or operation faster in SGX



# Evaluation

## Execution time

- TPC-H benchmarks, average of 5 runs, labeling from [Savvides et al.'17]



- HYDRA (PHE) on average 1.6x faster than Cuttlefish
- HYDRA (SGX) on average 11.3x faster than Opaque [Zheng et al.'17] (ORAM disabled)
- HYDRA (Hybrid) 1.7x and 1.6x times faster than HYDRA PHE and HYDRA SGX respectively
  - 2.7x and 17.9x faster than Cuttlefish and Opaque respectively

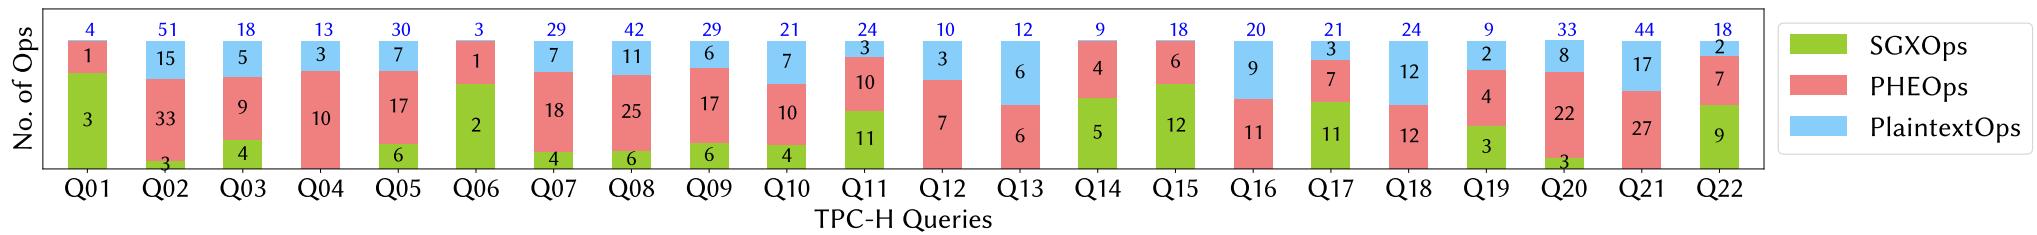
[Zheng et al.'17] W. Zheng, A. Dave,

J. Beekman, R. A. Popa, J. E. Gonzalez, and I. Stoica: Opaque: An Oblivious and Encrypted Distributed Analytics Platform. NSDI'17.

# Evaluation

## Transparency

- What if a programmer wanted to use/combine mechanisms manually?



- Proxy for effort: #ops and distribution for HYDRA Hybrid in TPC-H
  - Large number of operators gives idea of all possible combinations!
  - 16 of 22 queries use mix of PHE and SGX
  - Note: 18 of 22 queries use plaintext operations

# Extensibility

1. Adding/changing data relations (schema changes): (re-)label following  $\mathcal{L}$  (data manager w/ security expert)
2. Modifying security policy  $\mathbb{S}$ 
  - A. To assign mechanisms differently to levels: change only  $\mathbb{S}(l)$  (security expert)
  - B. To add/remove levels: change  $\langle \mathcal{L}, \preccurlyeq, \sqcup, \sqcap \rangle$  and  $\mathbb{S}(l)$  (security expert)
    - i. Possibly revisit data labeling (data manager w/ security expert, cf. 1.)
3. Adding whole new mechanism to  $\mathcal{D}$  or  $\mathcal{S}$ 
  - Some glue code
  - Augment heuristic (security expert w/ performance engineer)
  - Anybody wanting to use either 2.A. or 2.B.

# Ongoing/Future Work and Open Challenges

## HYDRA

- Placement
- Heuristic framework
  - More advanced heuristics
  - Improved verification of properties
  - Tracing back compilation/execution errors
- Security
  - Stronger guarantees, e.g., intermediate results, relation size leakage
  - Precision of guarantees, cf. universal composability [Canetti'00]
  - Beyond HbC: integrity

[Canetti'00] R. Canetti. Universally Composable Security: A New Paradigm for Cryptographic Protocols.  
<https://eprint.iacr.org/2000/067>.

# Ongoing/Future Work and Open Challenges

## Beyond HYDRA

- Multi-party computing (MPC)
- Differential privacy (DP) [Dwork et al.'17]
  - [Dwork et al.'17] C. Dwork, F. McSherry, K. Nissim, and A. Smith: Calibrating Noise to Sensitivity in Private Data Analysis. Journal of Privacy and Confidentiality 7(3) 2017.
  - [Pettai&Laud'15] M. Pettai and P. Laud: Combining Differential Privacy and Secure Multiparty Computation. ACSAC 2015.
- Continuous processing
  - Spark allows incremental computing, micro-batching
  - Main efforts in general still around data stores, e.g., *data clean rooms*
  - Vs MPC and/or DP [Pettai&Laud'15][Cao et al.'19]

# Corporate and Open Community Initiatives

- Fortanix: specific products and generic platform, PHE & SGX (\$130M raised)
- Opaque Systems: *data clean rooms* based on Opaque (\$32M raised)
- Decentrique: *data clean rooms*, SGX (\$26M raised)
- Decision3: chain(?) and SGX & AWS (\$100K raised)
- <Your company here>
- Veracruz: backed by Linux Foundation's Confidential Compute Consortium (C3)
- Apache Teclave: generic secure computing with SGX & Trustzone
- ...
- General lack: continuous processing support, security guarantees that are verified ("securified")



[Crunchbase]

# Thank You